

# Evaluating Security Hardening Options in GCC

Jacob Inwald



4th Year Project Report  
Computer Science  
School of Informatics  
University of Edinburgh

2024

# Abstract

NOT STARTED, need to make methodology concrete first

# Research Ethics Approval

This project was planned in accordance with the Informatics Research Ethics policy. It did not involve any aspects that required approval from the Informatics Research Ethics committee.

## Declaration

I declare that this thesis was composed by myself, that the work contained herein is my own except where explicitly stated otherwise in the text, and that this work has not been submitted for any other degree or professional qualification except as specified.

(Jacob Inwald)

# Acknowledgements

Any acknowledgements go here.

# Table of Contents

1	Background	1
1.1	Static Analysis . . . . .	1
1.2	Fuzz Testing . . . . .	2
1.3	Security Hardening . . . . .	2
1.4	Leading Examples . . . . .	3
1.4.1	Stack Canaries . . . . .	3
1.4.2	Fortify Source . . . . .	4
2	Methodology	6
	Bibliography	7

# Chapter 1

## Background

C and C++ are both prone to a class of bugs known as memory safety errors (memory errors). Memory errors occur when memory is accessed in an undefined manner, allowing the program to write/read in unintended ways to arbitrary regions of memory. This can lead to critical vulnerabilities that attackers can exploit. These errors have long been known about and long been exploited, from the Morris Worm in 1989 (Steeley [1989]) to Heartbleed in 2012 (Sass [2015]). Indeed, Microsoft has found that 70% of all its security defects in 2006-2018 were memory safety failures (Cimpanu [2019]), and the Chrome team similarly found 70% of all its vulnerabilities are memory safety issues (Cimpanu [2020]). Programming languages have since been developed that are memory safe and prevent such errors, but C and C++ remain widespread due to their speed and precision. It is infeasible to rewrite the language to prevent memory errors, and migrating away from C and C++ is not always possible for some hardware. Therefore, other approaches are required to aid in the prevention of memory errors. The key ones to discuss sit at development level (static analysis and fuzz testing) and production level (compiler hardening options).

### 1.1 Static Analysis

Static analysis for security assurance has been around since ITS4's release in early 2000; a simplistic syntactic matcher to rules that indicated vulnerabilities i.e. use of `strcpy()`?. Since then, static analysis tools have become significantly more sophisticated and complex, but the aim has always been the same - to catch security problems without executing the code. Static analysis tools will examine the program source code for flaws, marking sections of code that have potential bugs for a human programmer to then go over and resolve. These tools are incredibly effective at finding known bugs, but often rely on bugs being known. One of the more famous examples of this was Heartbleed, a memory error which was overlooked by static analysis tools until after its discovery and subsequent use in exploits (Sass [2015]). Another key challenge with static analysis is the reduction of false positives and false negatives. Static analysis tools walk the line between

false positives (reporting a bug where there is none) and false negatives (not reporting a bug where there is one). Both situations lead to different outcomes, with false positives leading to overhead for developers and false negatives leading to a false sense of security. Static analysis is now a strongly recommended part of the development lifecycle (SEI CERT C and C++), but is known to be fallible and is not a “cure-all” to security bugs.

## 1.2 Fuzz Testing

A fuzz tester (fuzzers) is a tool that iteratively mutates random input in an attempt to find security vulnerabilities for a piece of software. Fuzzers have been proven to be incredibly effective; as of 2023, a prominent fuzzer, OSS-Fuzz, has located over 10,000 vulnerabilities and 36,000 bugs across 1,000 projects (Google [2023]). Fuzzers differ to static analysis in that they execute the code while testing. Fuzzers rely on instrumentation, where extra code is added at compile-time to allow the fuzzer more access to the internals of program. This instrumentation can vary from sanitisers, which throw warnings or errors when dangerous behaviour is detected from the program i.e. out of bounds access, to tracing the execution path of the software or control flow (LLVM [2024]). The instrumentation will guide the mutations made by the fuzzer to iteratively improve vulnerability discovery. The upshot of this is that fuzzers can discover novel vulnerabilities not specified by the creator of the tool. These tools provide powerful security testing abilities, however, like static analysis, they rely on the developers using them effectively during development and the production lifecycle.

## 1.3 Security Hardening

It is almost guaranteed that production grade software will contain flaws. While static and dynamic approaches have been effective in culling these flaws, some will still slip through the gaps (i.e. Heartbleed). Therefore, security hardening is often used alongside these approaches. Security hardening attempts to prevent the memory unsafe behaviour of a language by introducing checks or transformations to workaround the potential flaws in the program. In GCC, different hardening features are enabled via the use of different command-line options such as `-fstack-protector-strong` or `-fPIe`. There are numerous hardening options to pick from, with each option aiming to solve a different problem. For example, the option `-fPIe` enables Address Space Layout Randomization, making return2libc attacks difficult on 32-bit systems and near impossible on 64-bit systems ?. Shown in Figure 1.1, is a list of security compiler options for GCC (and Clang and Binutils TODO: Remove those ones). One of those options, `-fhardened`, enables a subset of other compiler options to provide a default hardening configuration for a “naive” developer.

## 1.4 Leading Examples

In the previous section, the concept of security hardening was discussed, alongside some brief explanations of options available to a programmer. In this section, we will discuss the chosen flags for analysis, alongside some engineered examples to demonstrate some common use cases that they attempt to solve. We will start with the simpler flags, and then continue from there.

### 1.4.1 Stack Canaries

Stack canaries [Cowan et al., 1998] are a security hardening option that aim to prevent buffer overflows on the stack. The concept behind their implementation is very simple: some known word value is set before the return pointer on the stack frame and after the local variables, shown in Figure 1.2. This value is then checked against the known value before the program jumps to the return address, and aborts the program if the canary has changed. In order for the local variables to overwrite the return address in the stack frame, the canary will need to be overwritten. This means that an attacker needs to know the value of the canary in order to bypass the protection, which can prove challenging.

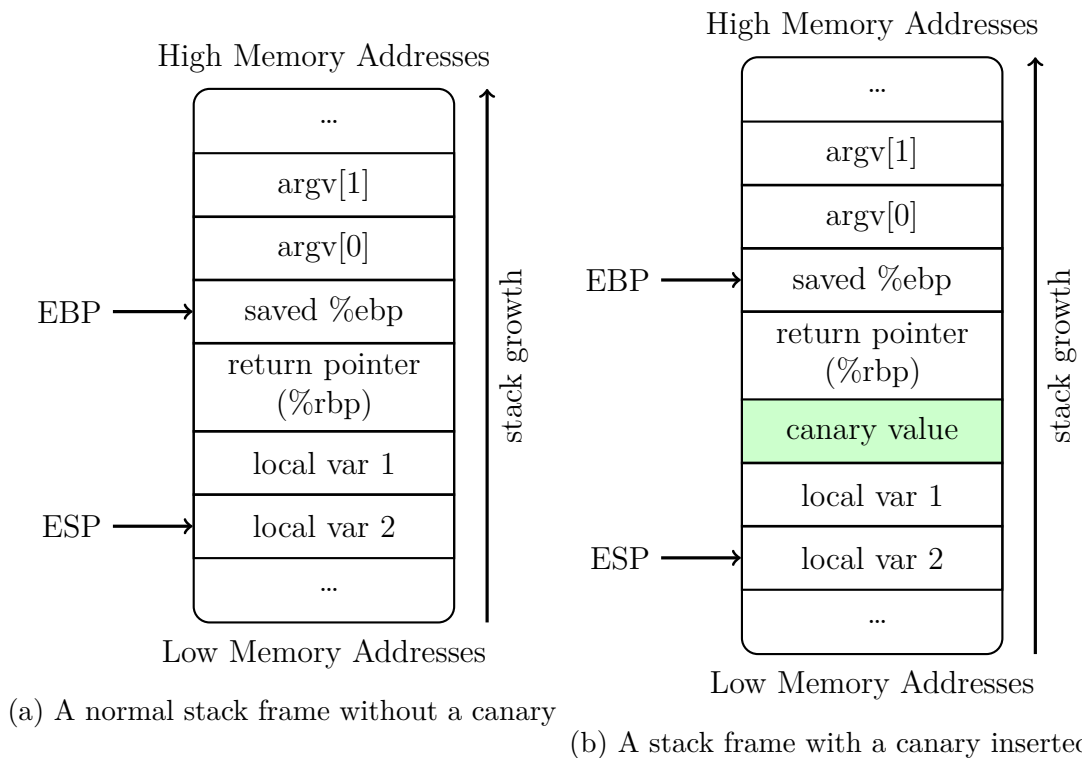


Figure 1.2: Stack frame with and without canaries

In Figure 1.3, we see some sample assembly code after stack canaries have been added in. As shown, the canary value is loaded in the word between the return address and then checked before the function returns to the return address given.



If the canary is different to what it was set to, the program throws a stack check fail exception and then exits.

```

1 ; start of main()
2 push  %rbp          ; preserves previous return pointer
3 mov   %rsp,%rbp     ; save return pointer to %rbp
4 sub   $0x40,%rsp     ; makes sure there's space for local variables (changes from 0x32
    ↪ to 0x40 with and without canaries)
5 mov   %edi,-0x34(%rbp)
6 mov   %rsi,-0x40(%rbp)
7 ; inserted code
8 mov   %fs:0x28,%rax ; load canary value into %rax
9 mov   %rax,-0x8(%rbp) ; store canary value 8 bytes before the return address
10 xor   %eax,%eax
11
12 ; start code block for main, left unchanged
13 ;
14 ; end code block for main
15
16 ; inserted code
17 mov   -0x8(%rbp),%rdx ; load canary value into %rdx
18 sub   %fs:0x28,%rdx   ; compare with saved canary value
19 je    0x4011c9 <main+147> ; skip error if not broken
20 call  0x401030 <__stack_chk_fail@plt> ; throw stack smashing error
21 ; unchanged
22 leave
23 ret   ; exit

```

Figure 1.3: x86 assembly code after compilation with the flag

Stack canaries are effective in preventing overflow attacks on the return address, where the buffer is overflowed sequentially from a local buffer. However, this doesn't prevent corruption of local variables. For example, in Figure 1.2, if local var 1 contains a boolean value used in a check and local var 2 can be arbitrarily written, then an attacker can change the value of local var 1 with impunity. This could lead to them passing checks they shouldn't be able to. Regardless, canaries are often a simple and relatively effective method of preventing sequential buffer overflows affecting the return address of a stack frame.

### 1.4.2 Fortify Source

The

Compiler Flag	Description
<code>-D_FORTIFY_SOURCE</code>	Fortify sources with compile- and run-time checks for unsafe libc usage and buffer overflows. Some fortification levels can impact performance. Requires <code>-O1</code> or higher, may require prepending <code>-U_FORTIFY_SOURCE</code> .
<code>-D_GLIBCXX_ASSERTIONS</code>	Precondition checks for C++ standard library calls. Can impact performance.
<code>-fstrict-flex-arrays=3</code>	Consider a trailing array in a struct as a flexible array if declared as <code>[]</code>
<code>-fstack-clash-protection</code>	Enable run-time checks for variable-size stack allocation validity. Can impact performance.
<code>-fstack-protector-strong</code>	Enable run-time checks for stack-based buffer overflows. Can impact performance.
<code>-fcf-protection=full</code>	Enable control-flow protection against return-oriented programming (ROP) and jump-oriented programming (JOP) attacks on x86_64
<code>-mbranch-protection=standard</code>	Enable branch protection against ROP and JOP attacks on AArch64
<code>-Wl,-z,nodlopen</code>	Restrict <code>dlopen(3)</code> calls to shared objects
<code>-Wl,-z,noexecstack</code>	Enable data execution prevention by marking stack memory as non-executable
<code>-Wl,-z,relro</code> or <code>-Wl,-z,now</code>	Mark relocation table entries resolved at load-time as read-only. <code>-Wl,-z,now</code> can impact startup performance.
<code>-fPIE -pie</code>	Build as position-independent executable. Can impact performance on 32-bit architectures.
<code>-fPIC -shared</code>	Build as position-independent code. Can impact performance on 32-bit architectures.
<code>-fno-delete-null-pointer-checks</code>	Force retention of null pointer checks
<code>-fno-strict-overflow</code>	Integer overflow may occur
<code>-fno-strict-aliasing</code>	Do not assume strict aliasing
<code>-ftrivial-auto-var-init</code>	Perform trivial auto variable initialization
<code>-fexceptions</code>	Enable exception propagation to harden multi-threaded C code
<code>-fhardened</code>	Enable pre-determined set of hardening options in GCC

Figure 1.1: From the Compiler Options Hardening Guide for C and C++

# Chapter 2

## Methodology

In this section we discuss our approach to collecting and interrogating data. We are aiming to investigate the real-life functionality

# Bibliography

- Catalin Cimpanu. Microsoft: 70 percent of all security bugs are memory safety issues, 2019. URL <https://www.zdnet.com/article/microsoft-70-percent-of-all-security-bugs-are-memory-safety-issues/>.
- Catalin Cimpanu. Chrome: 70 percent of all security bugs are memory safety issues, 2020. URL <https://www.zdnet.com/article/chrome-70-of-all-security-bugs-are-memory-safety-issues/>.
- Crispan Cowan, Calton Pu, Dave Maier, Jonathan Walpole, Peat Bakke, Steve Beattie, Aaron Grier, Perry Wagle, Qian Zhang, and Heather Hinton. Stack-guard: automatic adaptive detection and prevention of buffer-overflow attacks. In USENIX security symposium, volume 98, pages 63–78. San Antonio, TX, 1998.
- Google. Oss-fuzz, 2023. URL <https://google.github.io/oss-fuzz/>.
- LLVM. Sanitizer coverage, 2024. URL <https://clang.llvm.org/docs/SanitizerCoverage.html#tracing-control-flow>.
- Jeff Sass. The role of static analysis in heartbleed. 2015. URL <https://www.giac.org/paper/gsec/36189/role-static-analysis-heartbleed/143117>.
- Don Steeley. A tour of the worm. 1989. URL <https://collections.lib.utah.edu/details?id=702918>.